

Second Look

A two-minute test of how well a volunteer sees a creek, and a model that may only ask

Alex Velazquez and Rachel Selbrede, September 2026

Live: <https://second-look-79t.pages.dev> · Code: <https://github.com/alejandro-publius/second-look>

Abstract

Citizen observations of creeks are uneven. People notice smell, foam and colour and walk past concrete banks, dug-out channels, plants that do not belong and pipes. Second Look measures that, per person and per feature, with a two-minute test on 16 real photos, and stores the score with every observation the person later makes, in OneAquaHealth's own FHIR profiles. Vision models take the same test. A model may only ask a volunteer to look again, on a feature it passed, after the person has answered; code, never the model, writes the record. Four Claude models took the test three times each. Every model passed built banks, and no model passed plants that do not belong, where almost every answer was "can't tell": the plant photos show no water, and the instruction says to answer can't tell when no stream is in view. On 46 frames of open creek footage the gate dropped 29 of 64 candidate flags, each on a feature that model had not passed. The HL7 validator found 0 errors in 17 records checked against OneAquaHealth's guide. Neither wave retained an eligible completed sitting, so we could not estimate a human benefit. The human rows in section 3.1 report the pre-registered analyses and link to their exclusions. The late retry of the second lock kept the original cutoff; docs/deviations.md records the delay.

1. The problem

People judge a creek the way they judge a park. Tidy and green reads as healthy. OneAquaHealth's project lead said it in the first workshop: volunteers catch smell, foam and colour, and walk past concrete banks, a channel that was dug out, and pretty plants that do not belong. So the best-looking creek can get the best rating and deserve the worst.

Professional surveyors fixed this long ago. In the UK's River Habitat Survey, "only surveys from accredited surveyors will be entered on the RHS database", and accreditation means a course and a test (RHS manual 2003, pages 3 and 20; docs/notes/sources.md). Some volunteer programs certify people for a method, such as water chemistry. None we know of measures how well each volunteer sees each feature, or keeps that score with every observation, so a volunteer's observation arrives with no mark of how well its observer sees.

Measure each volunteer, per feature, and store the measure with the data. The analyst sees "4 of 4 on built banks, tested Sep 23" beside an answer, never a blended grade or a probability. Follow-up questions are chosen by code from the answers, the person's scores and the weather, two at most: "It has not rained here for N days. Is anything coming out of that pipe?" The AI takes the same test and earns the right to ask one question, feature by feature.

Every risk, what stops it, and the test that proves it

What can go wrong	What stops it, by construction	Proof
A volunteer walks past a built bank	The lesson teaches the four features people miss, and the test measures each one; the score is stored with every answer	<code>content/lessons/, core/scoring.py, core/tests/test_scoring.py</code>
Someone taps at random	Four items per feature, two present and two absent, so chance scores about 2 of 4; a low scorer who answers No is asked for a photo	<code>content/test_items.yaml, core/followups.py</code> rule <code>low_score</code>
A pretty creek gets the best rating	A good rating beside a reported built bank, sewage or plant that does not belong triggers one question asking the person to keep or change it, in the creek check and in a video walk	<code>content/followups.yaml</code> rule <code>rating_check</code> , <code>core/tests/test_followups.py</code> , <code>core/tests/test_walks.py</code>
A pipe report means nothing without the weather	The dry pipe question fires only after dry days from Open-Meteo, and is skipped when the weather is unknown	<code>core/rainfall.py, core/tests/test_rainfall.py</code>
The model invents a feature	Model output becomes a Flag through the gate or is dropped; a fuzz test throws arbitrary output at it	<code>core/gate.py, core/tests/test_gate.py</code>

What can go wrong	What stops it, by construction	Proof
The model was never good at that feature	A model may flag only a feature it passed on the same test as the people; the pass table is a committed file, and one from the fake client licenses nothing	<code>results/model_pass_table.json</code> , <code>core/tests/test_checker.py</code>
Duplicate and test pins fill the map	A precise pin within 30 metres of an existing spot is offered as that spot; names that read like a test are kept out of the counts and listed apart for a person to check; a coarse pin is never compared	<code>core/act.py</code> , <code>core/tests/test_act.py</code>
A record a city's systems cannot read	Sample records from both emitters, the Python API and the TypeScript Worker (17 files, walk records among them), are validated against their guide in CI, and golden vectors hold the live emitter to the same output	<code>scripts/fhir_validate.py</code> , <code>worker/test/golden.test.ts</code>
Someone edits history	A hash-chained audit log, checked by a script	<code>audit/</code> , <code>scripts/verify_audit.py</code>
Someone rewrites the audit log, last line included	Once a day its last hash is stamped with OpenTimestamps, if it has changed since the last stamp. A rewrite with fresh hashes still holds together, so the stamped line is compared with the log's own line, by the script and by <code>/verify</code> in your browser	<code>uv run python scripts/verify_audit.py</code> , then <code>uv run python scripts/ots_status.py</code> , which fails when the log no longer has a stamped line; <code>.venv/bin/ots verify proofs/audit-head-2026-09-24.ots</code> checks the stamp itself
The plan was written after the data came in	The <code>prereg-v1</code> tag object and <code>docs/analysis_plan.md</code> are stamped with OpenTimestamps, a public timestamp service that anchors hashes in Bitcoin; it is not our own chain. The <code>prereg-v2</code> and <code>prereg-v3</code> tags and their plans are stamped the same way, under <code>proofs/</code> . <code>/verify</code> shows each proof and its Bitcoin block once confirmed	<code>.venv/bin/ots verify proofs/prereg-v1.tag.ots</code> , <code>.venv/bin/ots verify -f docs/analysis_plan.md proofs/analysis_plan.md.ots</code> (<code>uv sync</code> puts <code>ots</code> in <code>.venv/bin</code> ; it needs a Bitcoin node to finish), or <code>uv run python scripts/ots_status.py</code> , which checks the block against a public explorer; <code>proofs/README.md</code>
We fool ourselves with the statistics	The analysis plan is tagged before any data; every README number is checked against <code>results/</code> in CI; a synthetic result is cited only where the sentence says it is a simulation	<code>docs/analysis_plan.md</code> at <code>prereg-v1</code> , <code>scripts/verify_claims.py</code>
Judge mode leaks the answer key	Before the lock, this: judge mode is shut by a lock constant while a wave of the study runs (until Sep 28 for the first wave, until the second lock for the second), and its answer routes refuse before then on the Worker and in the Python API. After the lock, nothing: judge mode says right or wrong on the same 16 photos the test uses, so anyone can rebuild the key and carry a perfect score into their creek checks, and a retake in a fresh browser does the same. Known weaknesses says so, and that it was open for two days between the waves	<code>core/lock.py</code> , <code>apps/api/tests/test_study.py::test_demo_answer_is_shut_before_the_lock</code> , <code>worker/test/e2e.mjs</code> ; what is not stopped: <code>docs/THREAT_MODEL.md</code>
A refresh loses a session	The session resumes from the server; the creek check queues offline and sends later	<code>apps/web/lib/offline.ts</code> , <code>apps/web/tests/</code>

2. Method

How a model's answer becomes a flag, or nothing, is drawn in `docs/diagrams/ai-gate.svg`, too wide to print here; 2.2 gives the same steps in words.

2.1 The test

Sixteen photos, four per feature, two with the feature and two without, in a new random order for each person. The answers are Yes, No and Can't tell, and Can't tell counts as wrong. The score is kept per feature, out of 4, in the record of the test sitting; a dated qualification and a Provenance link carry it to every Observation the person makes. The question wording and the analysis were written into `docs/analysis_plan.md` and tagged `prereg-v1`, and the key was frozen as a hash in `results/key_hash.json`, on Sep 21, the day the test opened, before the first paid model run on Sep 23. The dates here are Pacific time.

2.2 The gate

A vision model can help a volunteer look again. It can never decide what is stored. Every model answer takes this path:

1. **The person answers first.** The creek check asks the official app's questions in its order (`content/form.yaml`), in the app's own words, quoted from its public translation bundle and checked again by `make app-strings-check`; the answers are the person's own.
2. **The model is asked only where it passed.** `core/checker.py` reads the committed pass table, `results/model_pass_table.json`, and does not even ask about a feature that model did not pass on the same 16-photo test people take. A table that does not say `"real": true` licenses nothing.
3. **Its answer is forced.** One photo, one feature, the frozen question wording. Anything that is not yes, no or can't tell with a short note becomes can't tell and counts as malformed (`force_answer` in `core/checker.py`).
4. **The gate turns it into a flag or drops it.** `parse_flags` in `core/gate.py` keeps a `Flag` only for a known feature the model passed, with a sane confidence and a short plain note, and drops everything else with a plain reason. It never raises. On the footage run it dropped 29 of 64 candidate flags, each for a feature that model had not passed.
5. **A flag can only make one question eligible.** `core/followups.py` picks the follow-up questions from the answers, the rain (`core/rainfall.py`), the person's scores and the flags, by the rules in `content/followups.yaml`: two at most, the model's at most one, no model call inside.
6. **The person taps "I looked again" or "Skip", and no stored answer changes.** The question comes only for a feature the check asks about, and the model's note is shown only as "the checker noticed" (`core/followups.py`, `docs/MODEL_CARD.md`).
7. **The record is built from human inputs only.** `build_record` in `core/gate.py` has no parameter that could carry a flag, a model id or model text. `core/fhir_emit.py` writes the record under OneAquaHealth's profiles with the person's score, and `scripts/fhir_validate.py` checks it in CI.

Where a person meets it today: in part 2 of the test, offered after the score, and nowhere else a person sees. The assisted half of part 2 sees that one question, "The checker noticed something here. Look again?", from flags made once from the models' stored answers through the same gate (`results/assist_flags.json`); the note is not shown, and no model is called while anyone answers. Its judge mode, `/t2/demo`, shows the same from the second lock; until then `/how-we-know` shows the question a kept flag makes on one footage frame. The creek check and the walks do not show it yet. `scripts/build_walks.py` sends the footage run's answers through the same gate when the video walks are built, and no flag on those clips passed it, so 0 of the 3 walks shows a checker question. The live creek check runs with the checker off (`CHECKER_ENABLED`), so both servers pass it no flags and no model is called. `examples/footage-flag/` shows it act on the footage run.

2.3 What the AI cannot do

It cannot write the record. It cannot speak on a feature it did not pass, or on a made-up pass table. It cannot ask more than one question, or ask before the person answers. The gate says how, and Three properties that follow names each test. Two more limits:

It cannot	Enforced by	Test
Put more than a short labelled note in front of a person	The gate drops a note with markup, line breaks or text direction controls, and the page shows the note only after "the checker noticed", cut to 160 characters	<code>core/tests/test_gate.py::test_note_with_markup_or_a_direction_control_is_dropped</code> , <code>core/tests/test_checker.py::test_long_note_is_cut_to_160_and_flat</code>
State a risk for a named site	Every health or ecology sentence comes from <code>content/approved_sentences.yaml</code> with a source	<code>core/tests/test_healthcard.py</code> , <code>core/tests/test_act.py</code>

2.4 Three properties that follow

Property	Why it holds	The test that fails if it breaks
The model cannot write the record	The only way to make a record takes no model output	<code>core/tests/test_gate.py::test_build_record_signature_carries_human_inputs_only</code> , and the fuzz test <code>core/tests/test_gate.py::test_fuzz_model_output_never_reaches_answers_or_labels</code>
A model speaks only where it passed, and a made-up pass table licenses nothing	The gate and the checker both read the committed table and require <code>"real": true</code>	<code>core/tests/test_harden_gate_properties.py::test_every_kept_flag_is_licensed_for_that_exact_model</code> , <code>core/tests/test_checker.py::test_synthetic_pass_table_never_licenses_a_flag_by_default</code>

Property	Why it holds	The test that fails if it breaks
Code chooses the questions: two at most, the model's at most one, no network and no model call	Follow-up selection is a pure function	<code>core/tests/test_harden_followups_properties.py::test_the_content_table_never_asks_more_than_two_questions_for_any_input, core/tests/test_harden_followups_properties.py::test_the_selector_runs_with_http_and_the_model_client_patched_to_raise</code>

The hard parts of building this, and how each is proved, are in `WRITEUP.md`. The decisions are in `docs/adr/`.

3. Evidence

All numbers below come from files in `results/`, written by the scripts in `evals/` and checked against the text by `scripts/verify_claims.py` in CI. Failed runs stay in `results/` too.

3.1 The models on the test and on footage

	The 16-photo test people take	Frames from open creek footage
Pool	16 photos, 4 per feature, the frozen question wording	46 frames from 5 openly licensed videos in 3 countries, screened for people and text
Right answers, three runs of the 16 photos	Claude Haiku 4.5 31 of 48, Claude Sonnet 5 33 of 48, Claude Opus 5.5 33 of 48, Claude Fable 5.1 34 of 48; per feature, with intervals, in <code>results/benchmark_20260924T054939Z.json</code>	not measured: none of the frames has a label, so this column reports agreement
A checker that always answers No, the floor to beat	24 of 48	not asked
Right answers without the plant photos	Claude Haiku 4.5 31 of 36, Claude Sonnet 5 33 of 36, Claude Opus 5.5 33 of 36, Claude Fable 5.1 34 of 36	not asked
Which features each model passed	the table below	not asked: footage decides nothing, a pass is earned on the test
Agreement between models	not asked	Haiku 4.5 and Fable 5.1, the pair that agreed least, gave the same answer on 41 of 46 frames for built banks, 15 for a dug-out channel, 38 for invasive plants and 30 for pipes
Flags the gate stopped	not asked: an answer on the test is scored, never flagged	29 of 64 candidate flags dropped, 35 kept, because a model may flag only a feature it passed
One checker question about one frame, in cents	not asked	Claude Haiku 4.5 0.2, Claude Sonnet 5 0.4, Claude Opus 5.5 1.0, Claude Fable 5.1 2.3
Cost per 100 frames, each asked the four questions three times by all four models, with the run's adversarial frames counted in	not asked	51.3 USD, with direct calls at the full price
What a phone downloads in the background on a first visit, so a test started online keeps going offline and the creek check works offline after one visit	2.3 MB in 59 files, under the budget of 3 MB: the pages of the test and the creek check, and a phone-size copy of each of the 38 photos the test shows: the warm-up pair, the lesson and practice photos, and the 16 test photos	not asked

Which features each model passed on the 16-photo test: all four photos of a feature right in at least two of three runs (`results/model_pass_table.json`).

Model	Built bank	Dug-out channel	Invasive plant	Pipe running
Claude Haiku 4.5	passed	passed	did not pass	did not pass
Claude Sonnet 5	passed	passed	did not pass	did not pass
Claude Opus 5.5	passed	passed	did not pass	passed
Claude Fable 5.1	passed	did not pass	did not pass	passed

The exact model ids, as sent to the API in the run of 2026-09-24: `claude-haiku-4-5-20251001`, `claude-sonnet-5`, `claude-opus-5-5` and `claude-fable-5-1`. Their prices and the day each was checked are in `docs/notes/model_ids.md`.

The full loop, from a desk: 3 video walks from 3 countries, each ending in a FHIR record made on the phone. The HL7 validator checked 17 records against OneAquaHealth's guide, 3 of them walk records, with 0 errors.

No finished test from a person was kept before the data lock at 2026-09-28T01:00:00Z: 0 with the lesson and 0 without it, so there is no human row. The one pre-registered run, with what each of the plan's rules removed, is in [results/usability_20260929.md](#).

Does the checker's question help? Too few people finished part 2 for the plan's test: assisted 0, unassisted 0, and the plan needs 20 in each ([results/assist_20260929.md](#)). This measures one thing only: whether the checker's question helps a person whose first answer was wrong or Can't tell. It does not show that the checker cannot mislead anyone, because every flag in this set was correct (Known weaknesses).

The second wave, under plan prereg-v3: no finished test from a person was kept from 2026-09-30T04:00:00Z to the second lock at 2026-10-03T04:00:00Z: 0 with the lesson and 0 without it, so the second wave has no human row either. Its one run, with what each of the plan's rules removed, is in [results/usability_w2_20261004.md](#).

Does the checker's question help, in the second wave? Too few people finished part 2 for the plan's test: assisted 0, unassisted 0, and the plan needs 20 in each ([results/assist_w2_20261004.md](#)). This measures one thing only: whether the checker's question helps a person whose first answer was wrong or Can't tell. It does not show that the checker cannot mislead anyone, because every flag in this set was correct (Known weaknesses).

3.2 The benchmark and its size

The same 16 photos, 4 per feature (2 with the feature, 2 without), run 3 times more by `evals/benchmark.py`, a second set of calls from the ones that wrote the pass table. Each cell is right answers of all answers, with the Wilson 95 percent interval in percent. Can't tell counts as wrong. The file is [results/benchmark_20260924T054939Z.json](#).

Model	All 16 photos	Built bank	Dug-out channel	Invasive plant	Pipe running	Can't tell
Claude Haiku 4.5	31 of 48 (50 to 77)	12 of 12 (76 to 100)	10 of 12 (55 to 95)	0 of 12 (0 to 24)	9 of 12 (47 to 91)	35 percent
Claude Sonnet 5	33 of 48 (55 to 80)	12 of 12 (76 to 100)	12 of 12 (76 to 100)	0 of 12 (0 to 24)	9 of 12 (47 to 91)	27 percent
Claude Opus 5.5	33 of 48 (55 to 80)	12 of 12 (76 to 100)	9 of 12 (47 to 91)	0 of 12 (0 to 24)	12 of 12 (76 to 100)	31 percent
Claude Fable 5.1	34 of 48 (57 to 82)	12 of 12 (76 to 100)	10 of 12 (55 to 95)	0 of 12 (0 to 24)	12 of 12 (76 to 100)	27 percent

Twelve answers per feature is small. An interval such as 55 to 95 percent means the test cannot tell a good model from a very good one. Read the intervals, not the point numbers, and read them as too narrow. Each interval counts the twelve answers of a feature, or all the answers of the whole test, as independent. They are not: the three runs repeat the same four photos of a feature, and a model mostly answers a photo the same way each time (the per-answer sweep of the same night, [results/model_sweep_20260924T054756Z.json](#), shows it). Counted over the photos, each interval would be wider, so "each model beat the floor of always answering No" holds on the point numbers only.

This run kept counts, not replies. `make reproduce` checks that each accuracy and interval follows from the recorded counts, but the counts themselves, the Can't tell share and the count of malformed replies are as recorded: no reply is left to grade them from, and it prints a note saying so.

On creek footage. The four models also answered every feature question on 46 frames from 5 openly licensed videos in 3 countries. No frame has a label (0 labelled), so this measures agreement between models, not accuracy. The pair that agreed least, Haiku 4.5 and Fable 5.1, gave the same answer on 15 of 46 frames for a dug-out channel. Four frames drawn by code, a white one, a black one, a room made of flat rectangles and a screenshot of text (`evals/fixtures/__init__.py`), got can't tell as each model's majority answer over three runs, on every feature ([results/footage_latest.json](#), under `adversarial`; the run kept the majority answers, not each reply).

3.3 How the models fail

- **A run once measured our config instead of the models.** The first paid sweep sent an answer tool whose allowed values came from YAML, where an unquoted yes and no are booleans. The models were told to answer true or false, and 73 of 144 answers were malformed. The values are now quoted and pinned by `evals/tests/test_model_sweep.py::test_the_answer_tool_allows_exactly_the_three_words_as_strings`; the run done again had 0 malformed. Both runs stay in [results/](#).
- **No model passes invasive plants.** In the sweep behind the pass table, 47 of the 48 answers on the plant photos were can't tell. The plant photos are close views of a plant with no water in them, and the instruction says to answer can't tell when the photo does not show a stream. The notes show it: several name the plant, for example "possibly invasive Himalayan blackberry, but no stream is visible", and still answer can't tell. So this result measures our photos and our instruction as much as the models. The photos and the question wording are frozen with the tagged analysis plan, so changing them would be a deviation. The instruction is not in the plan: it is ours, in `evals/models.yaml`, and changing it would need a new paid run and a new pass table.

- **Can't tell is common.** Answers per feature in that sweep, all four models together:

Feature	Answers	Yes	No	Can't tell	Right
Built bank	48	24	24	0	48
Dug-out channel	48	20	24	4	42
Invasive plant	48	1	0	47	1
Pipe running	48	22	23	3	43

- **Models disagree on footage.** On the same frames, Haiku 4.5 answered yes for a dug-out channel on a share of 0.2 of them and Opus 5.5 on a share of 0.0. Without labels we cannot say which is right.
- **A pass can be luck.** Four photos, three runs. A pass on a feature says the model got these four photos right, not that it sees that feature in general.
- **The note is checked for form, not for truth.** The gate refuses markup, control characters and long notes, but a note that is wrong in plain words gets through. That is why a flag can only ask the person to look again, and the person's answer is what is stored.

3.4 Cost

Every paid call is one line in `results/cost_log.jsonl`, with its tokens and its price; the fake runs spend nothing. All paid calls together: 5113 calls for 41.6 USD, of which footage 34.2 USD, the sweeps 3.7 USD, the benchmarks 3.2 USD and the stored answers for part 2 0.5 USD. That includes the run that measured our config and the three-model run before Opus 5.5 and Fable 5.1 joined.

- The four-model run: the sweep 2.2 USD, the benchmark 2.2 USD and the footage 23.6 USD.
- Per 100 footage frames, four models, direct calls at the full price, with the run's adversarial frames counted in: 51.3 USD. The batch interface costs half as much, but a batch once waited three hours in the queue.
- The live site calls no model, so a volunteer's check, and a person's part 2, cost nothing in model calls.

4. The FHIR mapping

The records one visit makes, and how the score reaches every Observation, are drawn in `docs/diagrams/fhir-graph.svg`, too wide to print here; the list of what one visit produces, later in this section, names the same records.

OneAquaHealth says citizen data should stand beside lab data under the same profiles and value sets. A lab result is trusted because its quality checks travel with it. Second Look gives a volunteer's observation the same thing: their per-feature score, stored in the record of their test sitting beside a dated Practitioner qualification for the test, and linked through Provenance to every Observation they make. The four features are the ones the project lead named. The creek check asks the official Citizen Science App's questions in their order and in the app's own words and languages, quoted from its public translation bundle (`docs/notes/app_strings.md`). The city actions are OneAquaHealth's own restoration measures, from the [OneAquaHealth Policy Brief \(2026\)](#), page 9: replant margins, fix sewers, reconnect the floodplain, remove barriers, take out the concrete. The official app has no question for a dug-out channel, so the check asks none, and "reconnect the floodplain" waits for one; plants have no city measure of their own.

One creek visit becomes these resources (`docs/fhir_mapping.md`):

Resource	Profile	What it holds
Location, three nested	their Location0ah	creek, then reach (partOf creek), then spot (partOf reach, with position). Each has an identifier, a name and mode = instance, as the profile requires.
Practitioner	base R4	the volunteer, identified only by a random contributor token. One qualification coded second-look-test, period from the test date to 90 days later, issuer our Organization.
QuestionnaireResponse, test sitting	base R4	one item per feature with the score out of 4, computed by code. Authored by the Practitioner.
QuestionnaireResponse, visit	base R4	one item per answered check item. Link ids are the stable ids from <code>content/form.yaml</code> . Language is the language the questions were shown in, one of the official app's (UPDATE_32); the narrative is written in English and marked <code>lang="en"</code> , and says the language when it is not English, so the validator notes on those records that the two differ, which is true.
Observation, one per answered item	their ObservationIndicators0ah	status final, subject the spot, performer the Practitioner, effectiveDateTime, derivedFrom the visit response, a coded or quantity value.

Resource	Profile	What it holds
Provenance	base R4	target every Observation of the visit, agent author = Practitioner and assembler = our software Device, entity source = the visit response and the test sitting. This is the path from any answer to the score of the person who gave it.
Organization, Device	base R4	us, and the software.
Bundle, type collection	base R4	everything above. The emitter also produces a transaction Bundle with conditional creates for the sandbox mirror (to_transaction in core/fhir_emit.py).

- `Observation.value`: their TemporaryOahSystem#present and #absent. "Can't tell" is our second-Look#cant-tell because their system has no general unknown code.
- `Observation.category`: their indicator groups. Built banks and dug-out channel sit under #morphology (their spelling), invasive plants under #invasiveOrganisms, pipes and water height under #hydrology. The profile does not bind category, so this is additive.
- `Observation.code`: our feature codes artificial-bank, dug-out-channel, invasive-plant, pipe-running. Their profile binds code to OahIndicatorsNoHealthOahVs with strength preferred, so a code outside it produces an information note, not an error. The validator confirmed this. Water height uses their own #hydrology code directly.
- Coding displays must be the code system's own display. The words a person sees on screen live in `content/form.yaml`, never in the coding.
- Units are UCUM (m for water height).

One Digital Health and FAIR

The five dimensions and the FAIR principles, one line each

Story: a student walks to Strawberry Creek, takes a two-minute test on her phone, and from then on every observation she makes carries how well she sees each kind of damage. A city analyst reads her record beside a lab result under the same profile and knows how much weight to give each. The health card gives her one thing for herself, one for her dog and one for her city.

Five dimensions, in words, because no script produces them: citizen engagement, strong; education, strong; human and veterinary healthcare, partial (approved sentences for the person and the pet, no diagnosis, no site risk); industry 4.0, partial (FHIR records, a gated vision model, an audit log); environment, strong.

FAIR: findable through a Library entry in their sandbox and a public repository; accessible through a read-only FHIR endpoint; interoperable through their profiles, value sets and UCUM; reusable through MIT code, a pinned guide, Provenance on every record and a tagged analysis plan.

5. Limitations

- Part 2 cannot show that the checker never misleads. Every one of the 6 flags in its set points the way of the gold label (`results/assist_flags.json`), so the question only ever follows a first answer that was wrong or Can't tell. Part 2 measures whether the question helps that person, nothing more; a checker that is sometimes wrong would need wrong flags in the set.
- One labeller. The gold labels came from the picks file Alex wrote with the planner, a Claude chat (commit 81e62ed), and no second, blind label exists yet, so read every accuracy as agreement with this key (`docs/DATA_CARD.md`).
- Four photos per feature is coarse: enough to show a person what to practise and to flag an answer worth a second look, too coarse to weight votes by feature. In our simulation with made-up people (`results/consensus_coarseness.json`), weights from each feature's own photos did worse than a plain majority in all 5 skill patterns.
- A test started online keeps going if the network drops, from smaller copies of its photos. A reload while offline does not bring the sitting back: it opens the test at its start, or the offline page when the phone has not kept that exact link. The end of the lesson is not sent again once the network is back.
- The test and lesson photos are sent as the full-size JPEGs, so on a slow phone line each can take a few seconds to appear. That is on purpose: they are the study's pre-registered materials, and every person must see the same pixels; only the landing page and the poster serve smaller copies to everyone. A phone also keeps a smaller copy of each test and lesson photo, shown only when a photo cannot be fetched at all, mid-sitting.
- The official app has no question for a dug-out channel, so the creek check asks none: that score stands beside no answer, and the city measure it would lead to waits for one. Plants that do not belong have no city measure of their own.

- Of the 35 footage flags the gate kept, 32 are on a dug-out channel, so they ask nothing. The other 3 are one model's runs on built banks, all on one frame of creek water over stones seen from above, where the other models that passed built banks said can't tell. Passing four photos did not stop that flag, which is why a flag can only ask ([examples/footage-flag/](#)).
- The photos come from open collections in several countries and seasons, not from the creeks a Berkeley visitor will stand in.
- The footage labels come from the videos' own descriptions, and almost none names a feature, so the footage result leans on agreement between models, which is not accuracy.
- The citizen observer is modelled as a FHIR Practitioner, for want of a better fit in the guide; our proposal says so.
- The answer key can be rebuilt after the lock. Judge mode says right or wrong on the same 16 photos the test uses, and a retake in a fresh browser shows the score per feature each time, so a person who wants to can learn all 16 answers, take the test, keep 4 of 4 on every feature and carry that score into their creek checks, and nothing marks such a score. Judge mode is shut while a wave of the study runs, so it cannot leak the key to that wave, and the study counts only the first finished sitting from a browser. But it was open for two days between the waves, from the first lock on Sep 28 until Sep 29, and anyone who used it then may know the answers; it stores nothing, so nobody can say who did. It stayed shut until the second lock, Friday Oct 2, 2026 at 21:00 PDT, and reopened then ([docs/analysis_plan_v3.md](#), item 5). A separate bank of photos for judge mode and retakes would close it ([docs/THREAT_MODEL.md](#)).
- The second wave could not estimate a human benefit. Both arms were empty after the pre-registered exclusions. The results are published, and the failed job and its late retry are recorded in [docs/deviations.md](#).
- Human benefit remains unmeasured. Neither wave retained an eligible completed sitting. The first plan did not plan recruitment; plan v3 allowed a paid research panel alongside the public link. The published results report the exclusions and leave the human effect uncomputed. The study has had no ethics review; it was designed as an anonymous usability test that keeps no name or contact ([docs/analysis_plan_v3.md](#), [docs/deviations.md](#)).
- Their sandbox's name, `sandbox.hl7europe.eu`, stopped resolving on Sep 23, came back on Sep 28, and has come and gone since. `/two` shows their lab record from the last copy the Mac fetched and says when that was. The re-push of Sep 28 found 14 of our resources there and created one Observation that was not; their server refused the update of our Library entry with HTTP 400 ([fhir/sandbox_ledger.jsonl](#)).
- The dated and daily jobs run on one laptop, Alex's Mac, under his own logins: the copy of their record for `/two`, the iNaturalist line, the backup, the uptime check every 10 minutes, the OpenTimestamps stamp of the audit log, the sandbox re-push, the watch on our pull request and issues in their guide, and the one run of the analysis after each lock. While that laptop sleeps, they wait; the live site keeps working, and the lock itself is the Worker's clock, not a job ([DEPLOY.md](#), Jobs on the Mac).
- The tests step shows some tests as xfailed. Each is a known bug in our own code, in the gate, the FHIR writer, the referral check, the follow-up picker or the content loader, kept as a test of the right behaviour, marked as a strict expected failure so the run goes red the day it is fixed. Each is named in [docs/KNOWN_BUGS.md](#).
- The AI numbers come from paid calls on Sep 23, Pacific time: four models, three runs each, on 16 photos, once for the pass table and once more for the right-answer counts, plus the footage run. That is small. Read the intervals in the benchmark file, not the point numbers, and read them as too narrow: they count each run's answer as new, but the three runs repeat the same 16 photos and a model mostly answers a photo the same way each time ([results/model_sweep_20260924T054756Z.json](#) has every answer); counted over photos, each interval would be wider.
- Our own words are English only: the buttons, notes, follow-up questions, record pages, the test and its lessons. The creek check's questions and answers, and its Back, Next and Send, are in the official app's six languages, and a phone set to one of them opens the check in it; a translation whose meaning differs from the English shows in English, marked ([docs/notes/app_translations.md](#)). The app's Greek has no assessment questions, so it is not offered. A Spanish draft of our words waits for a fluent person to sign it.

6. Contributions

1. **A measured observer, carried in the record.** A two-minute test gives each volunteer a score per feature, and OneAquaHealth's own profiles carry it, through a qualification and Provenance, to every Observation that person makes; a city analyst reads each answer beside how well its observer sees.
2. **A model that may ask and never decide.** A model is licensed feature by feature by the same test people take, from a committed pass table; its output becomes a flag through one gate or is dropped, and the record is built from human inputs only.
3. **Measurement that shows its failures.** The pass table, the benchmark with its intervals and the plant photos that no model could judge are all in [results/](#), with the cost of every call.
4. **Open work, and work given back.** The code is MIT. What we contributed back, as the README lists it:

Sent to OneAquaHealth's implementation guide on 2026-09-24, in the open:

- [hl7-eu/oah pull request 5](#): our citizen observer example, five FSH files that build inside their guide with SUSHI 3.20.1 at no errors, and a proposal for carrying observer quality in their guide ([docs/ig_proposal.md](#)), with the gaps our validator runs found. It asks which resource should stand for a citizen observer.
- [hl7-eu/oah issue 6](#): their temporary code system spells one code morphology. We kept their spelling so our records validate.
- [hl7-eu/oah issue 7](#): `SpecimenOah.collection.collector` allows only a `PractitionerRole`, which leaves out a laboratory and a volunteer who takes a sample.
- [hl7-eu/oah issue 8](#): their sandbox's name stopped resolving on 2026-09-23, with the evidence from their own nameserver. It came back on Sep 28 and has come and gone since (Known weaknesses).

And for anyone, not sent to them: a read only MCP server over our own records, so any software agent can ask for a creek's records with the resource ids behind every answer ([examples/mcp/README.md](#)).

7. How to check this report

`make judge-check` runs the tests with no key and no network; `make verify-claims` checks every number in the README and the docs against `results/`; `make report-pdf` builds this file again from its sources, and a test fails when the PDF is older than them. The model card, the data card and the threat model are `docs/MODEL_CARD.md`, `docs/DATA_CARD.md` and `docs/THREAT_MODEL.md`. AI coding tools wrote most of the code and the text, from written briefs, and every change was checked by `make check` before it was committed; the people set the direction and made every decision that needs a person, as the README says under "How this was built".